

Introduction to ONNX

Open Neural Network Exchange

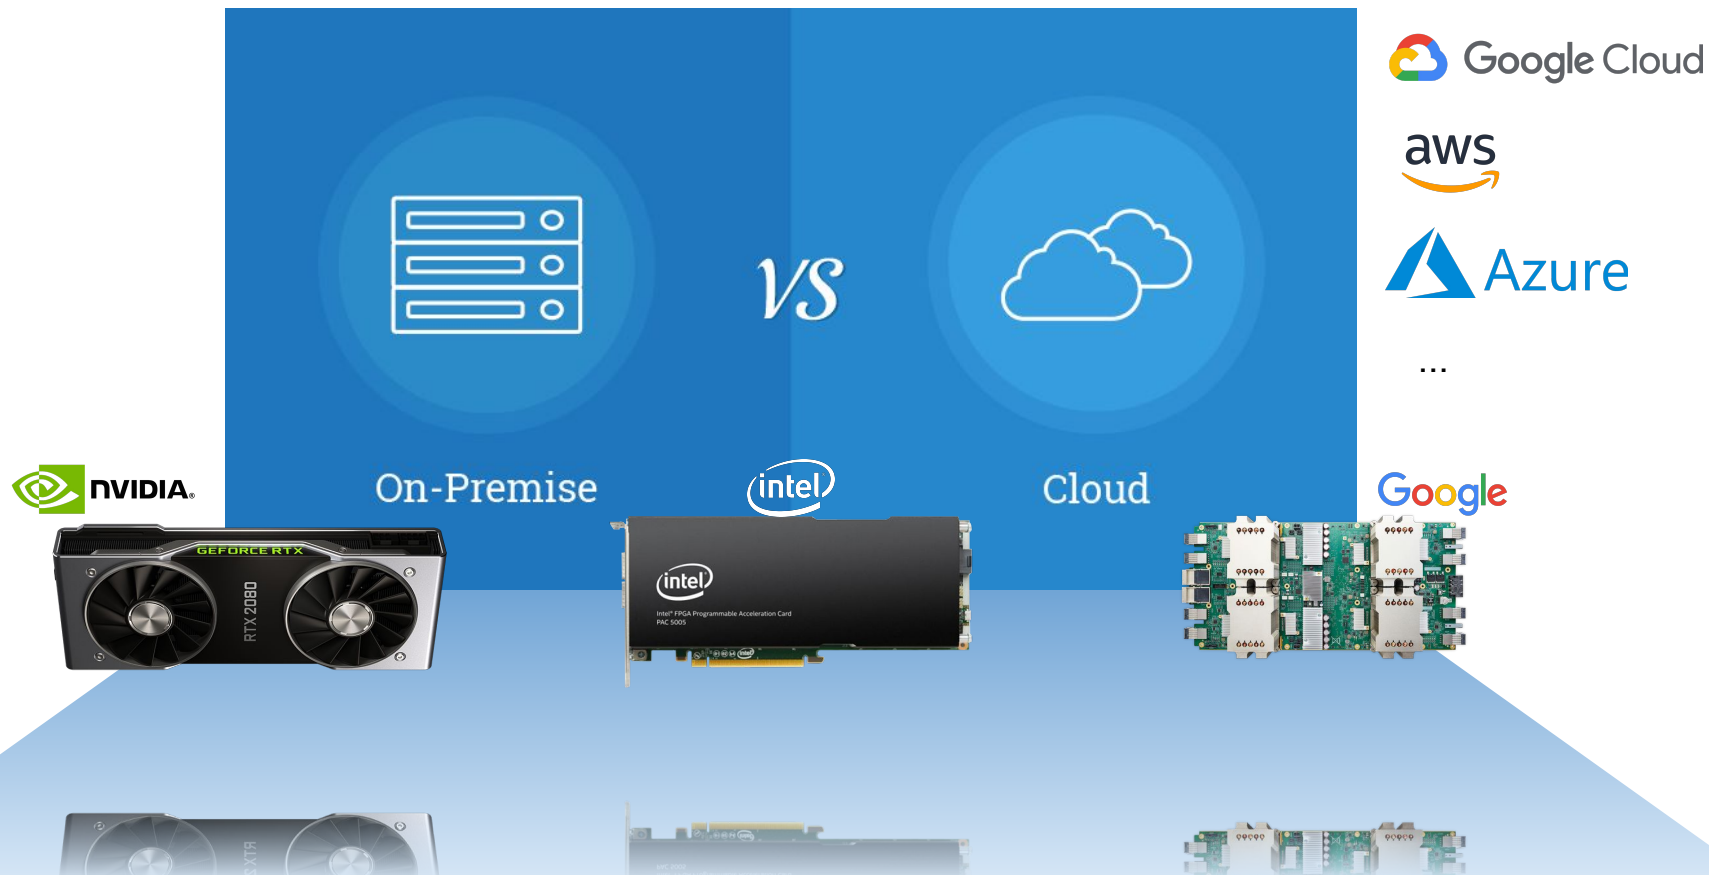
Objectives

- Challenges with Machine Learning;
- Open Neural Network Exchange
- ONNX Specification;
- ONNX Runtime;
- ONNX Ecosystem;
- ONNX Demos:
 - Create & use a new ONNX model;
 - Use an existing ONNX model;
- Where to find more information?






Challenges #1: ML Framework




Challenges #2: Select Training Infrastructure




Challenges #3: Select Inference Infrastructure



Challenges #4: Summary



Solution: Open Neural Network Exchange



ONNX Specification

MODEL PROPERTIES

format

ONNX v7

producer

skl2onnx 1.9.2

domain

ai.onnx

imports

ai.onnx v14

ai.onnx.ml v1

INPUTS

Age

name: **Age**

type: float32[, 1]

Fare

name: **Fare**

type: float32[, 1]

Pclass

name: **Pclass**

type: string[, 1]

Sex

name: **Sex**

type: string[, 1]

OUTPUTS

output_label

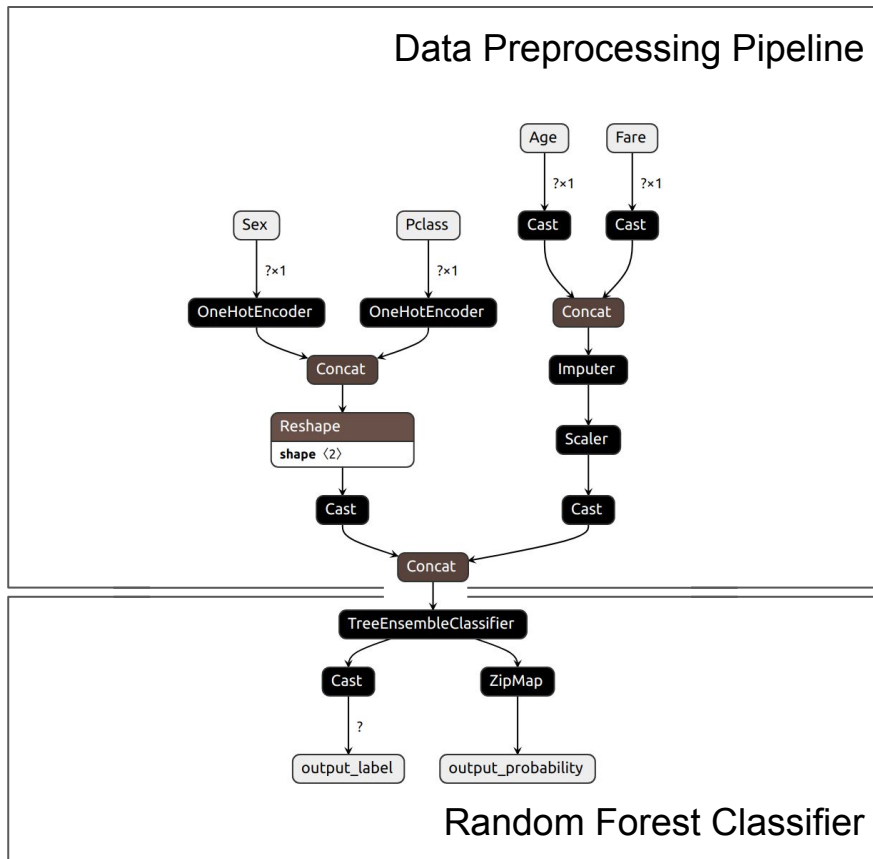
name: **output_label**

type: int64[]






output_probability

name: **output_probability**

type: sequence<map<int64, float32>>



ONNX Runtime



...

ONNX Ecosystem

ABBYY®

Alibaba Group
阿里巴巴集团

AMD

arm

aws

Baidu 百度

BECKHOFF

BITMAIN

cadence®

CEVA®

Facebook
Open Source

GRAPHCORE

habana

HAILO

Hewlett Packard
Enterprise

HUAWEI

IBM®

Idein Inc

intel® AI

MathWorks®

MAXAR

MEDIATEK

mi

Microsoft

NVIDIA.

NXP

OctoML

OPEN AI LAB
开放智能

Preferred
Networks

SIEMENS

SONY

Qualcomm

sas

商汤
sensetime

skymizer

SYNOPSYS®

Tencent

unity

verizon
media


vmware®

WOLFRAM

Yandex

ZETANE

Demo: Create & use a new ONNX model



Demo: Create & use a new ONNX model (#1)

```
from utils import create_preprocessor

from sklearn.ensemble import RandomForestClassifier
from sklearn.pipeline import Pipeline

# Creates a data preprocessing pipeline.
preprocessor = create_preprocessor(dataset)

# Defines the model as a pipeline which combines a data preprocessor and
# a classifier.
model = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', RandomForestClassifier(random_state=42))
])

# Trains the model.
model.fit(train.X, train.y)

# Uses the model for scoring.
model.predict(test.X)
```

Create SKLearn
Random Forest
Classifier

Demo: Create & use a new ONNX model (#2)

```
from utils import get_onnx_input_type
from skl2onnx import convert_sklearn

# Creates input type using dataset schema.
initial_type = get_onnx_input_type(train)
print(initial_type)

# Converts the model to the ONNX format.
onnx_model = convert_sklearn(model, initial_types=initial_type)

# Serializes the ONNX model to the file.
with open('model.onnx', "wb") as f:
    f.write(onnx_model.SerializeToString())
```

Convert and
Save ONNX
Model

```
[('Age', FloatTensorType(shape=[None, 1])),
 ('Fare', FloatTensorType(shape=[None, 1])),
 ('Pclass', StringTensorType(shape=[None, 1])),
 ('Sex', StringTensorType(shape=[None, 1]))]
```

} initial_type

Demo: Create & use a new ONNX model (#3)

```
from utils import get_onnx_input_data
import onnxruntime as rt

input_data = get_onnx_input_data(test)
sess = rt.InferenceSession('model.onnx')
pred, _ = sess.run(None, input_data)
```

Load and Use
ONNX Model

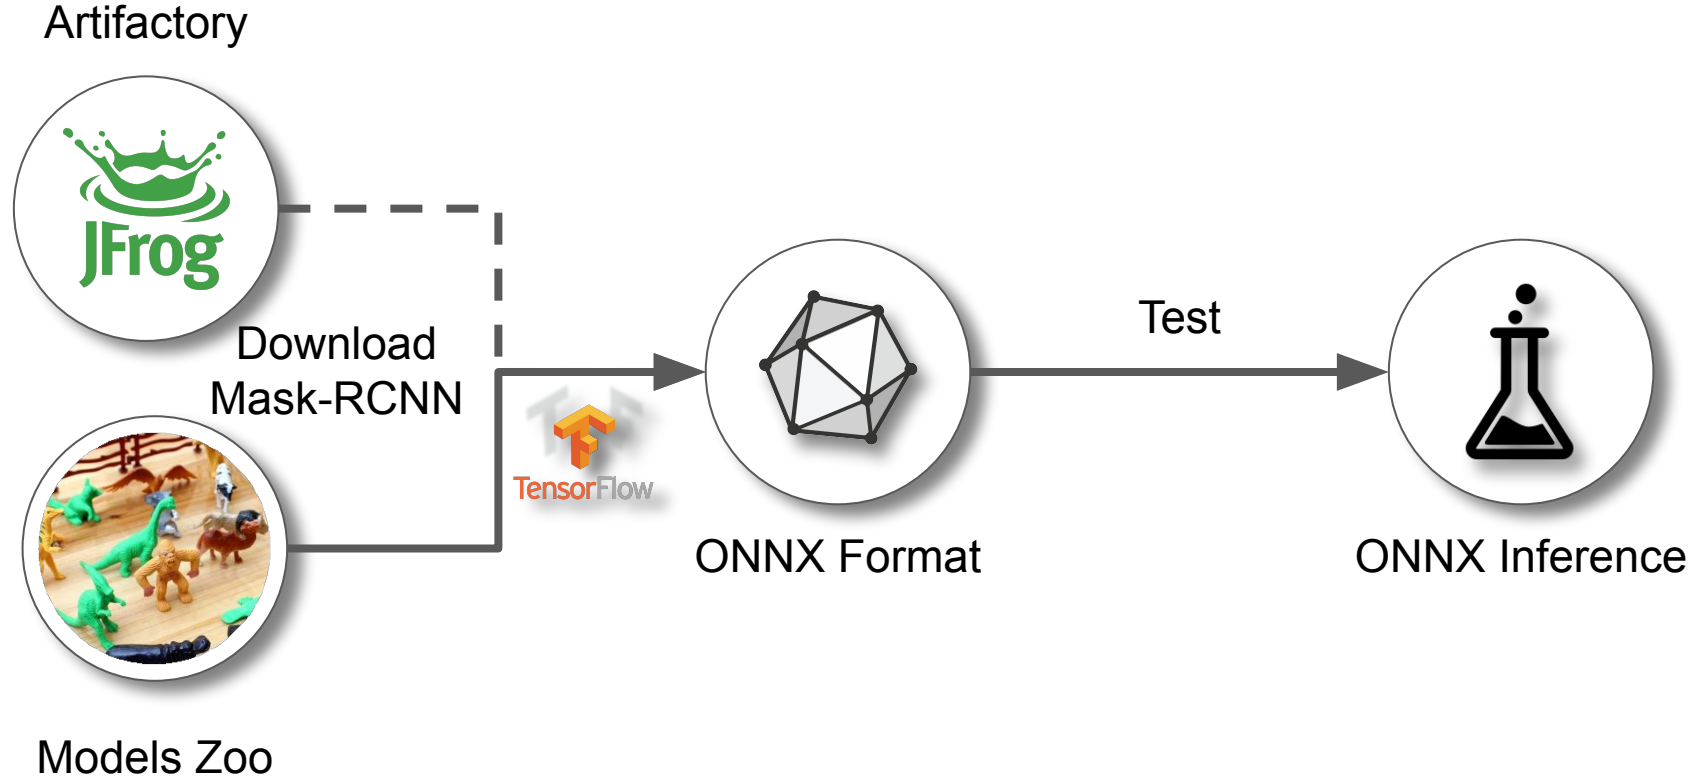
Input Data Transformation

	Age	Fare	Pclass	Sex
296	47.0	30.500000	1	male
682	25.0	41.579201	2	male
535	69.0	14.500000	3	male
644	56.0	35.500000	1	male
623	57.0	12.350000	2	male
..
377	19.0	7.775000	3	male
140	22.0	7.750000	3	female
173	56.0	30.695801	1	male
6	54.0	51.862499	1	male
731	28.5	16.100000	3	male



```
{'Age': array([
    [47. ], [25. ], [69. ], [56. ], 57. ], ..., dtype=float32),
 'Fare': array([
    [ 30.5  ], [ 41.5792], [ 14.5  ], [ 35.5  ], [ 12.35  ], ..., dtype=float32),
 'Pclass': array([
    ['1'], ['2'], ['3'], ['1'],  ['2'], ..., dtype=object),
 'Sex': array([
    ['male'], ['male'], ['male'], ['male'], ..., dtype=object])}
```

Demo: Use an existing ONNX model



Demo: Use an existing ONNX model (#1)

Get Pre-trained Mask-RCNN

```
!wget -O tmp/maskrcnn.onnx
```

```
https://github.com/onnx/models/raw/master/vision/object\_detection\_segmentation/mask-rcnn/model/MaskRCNN-10.onnx
```

Define Pre/Post Process Functions

```
def preprocess(image):
```

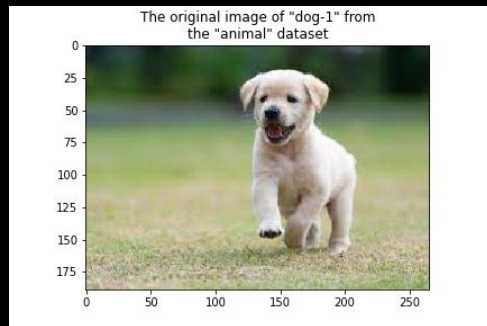
Convert the original image to Mask-RCNN input

```
def postprocess(image, classes, boxes, labels, scores, masks, score_threshold):
```

Take Mask-RCNN predictions and project to the original image...

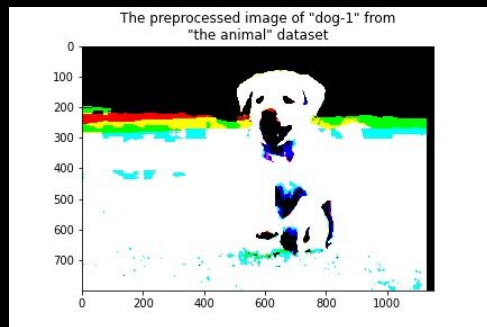
Demo: Use an existing ONNX model (#2)

```
orig_image = load_image(...)
```



Load The Original Image

```
pre_image = preprocess(orig_image)
```



Preprocess input for the Mask-RCNN input

Demo: Use an existing ONNX model (#3)

```
import onnxruntime as rt  
sess = rt.InferenceSession('maskrcnn.onnx')
```

```
f = open('coco_classes.txt')  
classes = f.read().split(',')
```

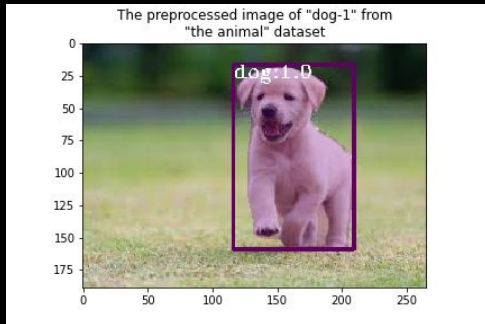
```
boxes, labels, scores, masks = sess.run(  
    None, {sess.get_inputs()[0].name: pre_image}  
)
```

```
post_image, _ = postprocess( orig_image, classes, boxes, labels, scores, masks, score_threshold)
```

Read ONNX model
and COCO classes

Make Prediction

Show Result



Where to find more information?

ONNX Home: <https://onnx.ai/>

ONNX GitHub: <https://github.com/onnx>

ONNX Model Zoo: <https://github.com/onnx/models>

ONNX Runtime: <https://onnxruntime.ai/>

Open Neural Network Exchange (ONNX) in the enterprise: how Microsoft scales ML: <https://www.youtube.com/watch?v=aHk7iUZDIk>

Netron UI: <https://netron.app/>

Netron GitHub: <https://github.com/lutzroeder/netron>

Thank You